

CS186 Discussion Section Week 5

SQL

Outline for Today

- Quick SQL Intro
- Except, Union, Intersect
- Nested Subqueries
- Joins
 - Inner, Left, Right, Full
- Aggregation

What is it?

- Structured query language
- AKA the most useful thing you will learn out of this course
- Used to communicate with databases
- ANSI standard for relational databases

Basic SQL

SELECT	target-list
FROM	relation-list
WHERE	qualification

- **target-list**: list of attributes in each relation (or *)
- **relation-list**: list of relations, usually a table
- **qualification**: set of select clauses

Tables

```
CREATE TABLE Students  
(sid text PRIMARY KEY,  
name text,  
gpa integer)
```

```
CREATE TABLE Courses  
(cid text PRIMARY KEY,  
name text,)
```

```
CREATE TABLE  
Enrollments  
(sid text,  
cid text,  
FOREIGN KEY(sid)  
REFERENCES students  
(sid),  
FOREIGN KEY(cid)  
REFERENCES courses  
(cid),  
PRIMARY KEY(sid, cid))
```

Basic SQL Examples

- `SELECT name FROM Students WHERE gpa = 3;`
- `SELECT * FROM Students;`
- `SELECT DISTINCT name FROM Students;`
- `SELECT name, cid
FROM Students, Enrollments
WHERE sid = sid;`

Basic SQL Examples

- `SELECT name FROM Students WHERE gpa = 3;`
- `SELECT * FROM Students;`
- `SELECT DISTINCT name FROM Students;`
- `SELECT name, cid
FROM Students, Enrollments
WHERE Students.sid = Enrollments.sid;`

Basic SQL Examples

- `SELECT name FROM Students WHERE gpa = 3;`
- `SELECT * FROM Students;`
- `SELECT DISTINCT name FROM Students;`
- `SELECT name, cid
FROM Students S, Enrollments E
WHERE S.sid = E.sid;`

Except, Union, Intersect

```
SELECT cid FROM Courses
```

```
EXCEPT
```

```
SELECT DISTINCT cid
```

```
FROM Students S, Enrollments E
```

```
WHERE S.sid = E.sid;
```

“Select all courses with no one enrolled.”

Notice that this is equivalent to subtraction.

Except, Union, Intersect

```
SELECT cid FROM Courses
```

```
EXCEPT
```

```
SELECT cid
```

```
FROM Students S, Enrollments E
```

```
WHERE S.sid = E.sid;
```

“Select all courses with no one enrolled.”

Notice that this is equivalent to subtraction.

Except, **Union**, Intersect

```
SELECT cid FROM Courses
```

```
UNION
```

```
SELECT cid
```

```
FROM Students S, Enrollments E
```

```
WHERE S.sid = E.sid;
```

“Select all courses.”

UNION removes duplicates, **UNION ALL** does not.

Except, Union, **Intersect**

```
SELECT cid FROM Courses
```

INTERSECT

```
SELECT cid
```

```
FROM Students S, Enrollments E
```

```
WHERE S.sid = E.sid;
```

“Select all courses with someone enrolled.”

INTERSECT removes duplicates, **INTERSECT ALL** does not.

Nested Subqueries

- **IN** operator: check if attribute is in subrelation
- **SELECT** name
FROM Students
WHERE sid **IN** (**SELECT** sid FROM **Enrollments**);
- “Select students enrolled in some class.”

Nested Subqueries

- **EXISTS** operator: check if subrelation is empty
- **SELECT** name
FROM Students
WHERE EXISTS (**SELECT** sid **FROM**
Enrollments);
- Wait, what?
 - EXISTS doesn't seem useful here...
 - Result: Correlated subqueries!
- This returns names of all students if the inner query finds one or more sid. If inner query does not return any sids, full query returns nothing.

Correlated Subqueries

- Inner query depends on outer query.
- SELECT name

FROM Students **S**

WHERE **EXISTS** (SELECT sid FROM Enrollments
WHERE sid = **S**.sid)

- Same query as before!
- Any Caveats?
 - Much slower. Inner query must run once per outer tuple.

Other Subquery Operations

- **UNIQUE**: is every result unique?
- **SELECT** name
FROM Students **S**
WHERE **UNIQUE** (**SELECT** sid FROM Enrollments
WHERE sid = **S.sid**)
- “Select students enrolled in exactly 1 class (since sid is a key of students).”

Other Subquery Operations

- $> ANY$, $> ALL$
- `SELECT name`
`FROM Students S`
`WHERE gpa \geq ALL`
`(SELECT gpa FROM Students)`
- “Select students with the highest GPA.”
- For the `WHERE` condition to be met, `gpa \geq` every single `gpa` returned by subquery
- In this case, subquery is not correlated.

Other Subquery Operations

- > ANY, > ALL
- SELECT name
FROM Students S
WHERE gpa = ANY
(SELECT gpa FROM Students)
- “Select all students!”
- For the WHERE condition to be met, gpa = any gpa returned by subquery
- Again, subquery is not correlated.

Summary of Operations

Operation	Truth condition
attr IN (subquery)	attr is one of the returned rows
EXISTS (subquery)	subquery returned non-empty result
NOT EXISTS (subquery)	subquery returned empty result
UNIQUE (subquery)	subquery returned all unique results
attr <i>op</i> ANY (subquery)	attr <i>op</i> (each returned row) returns true at least once
attr <i>op</i> ALL (subquery)	attr <i>op</i> (each returned row) returns true for all rows

Special Joins

```
SELECT (column_list)
FROM table_name
  [INNER | {LEFT | RIGHT | FULL } OUTER] JOIN table_name
  ON qualification_list
WHERE ...
```

- SELECT * FROM Students S
 INNER JOIN Enrollments E
 ON S.sid = E.sid;
- Not really different than before:
- SELECT * FROM Students S, Enrollments E
 WHERE S.sid = E.sid;

Outer Joins (LEFT)

- Outer joins add in unmatched rows.
- SELECT sname, cid FROM Students S

LEFT OUTER JOIN Enrollments E

ON S.sid = E.sid;

- Example output
- Rows of **left** relation that couldn't be matched are still here with a **NULL** **right** hand side.

sname	cid
Dan	160
Evan	186
Lu	186
Victor	186
Liwen	NULL
Mike	NULL

Outer Joins (RIGHT)

- Outer joins add in unmatched rows.
- `SELECT sname, cid FROM Students S
RIGHT OUTER JOIN Enrollments E
ON S.sid = E.sid;`
- Example output
- Rows of **right** relation that couldn't be matched are still here with a **NULL left** hand side.

sname	cid
Dan	160
Evan	186
Lu	186
Victor	186
NULL	188
NULL	189

Outer Joins (FULL)

- Outer joins add in unmatched rows.
- `SELECT sname, cid FROM Students S`

`FULL OUTER JOIN Enrollments E`

`ON S.sid = E.sid;`

- Example output
- Any unmatched rows are still there, with the other side **NULL**

sname	cid
Dan	160
Evan	186
Lu	186
Victor	186
NULL	188
NULL	189
Liwen	NULL
Mike	NULL

Aggregation

```
SELECT COUNT(*) FROM Students;
```

```
SELECT AVG(gpa) FROM Students;
```

```
SELECT COUNT(DISTINCT name) FROM  
Students;
```

```
SELECT DISTINCT COUNT(name) FROM  
Students;
```

count
12

GROUP BY

- GROUP BY aggregates rows together.

- SELECT standing
FROM Students
GROUP BY standing;

standing
freshman
senior
junior
sophomore

- SELECT standing, COUNT(*)
FROM Students
GROUP BY standing;

standing	count
freshman	122
senior	178
junior	212
sophomore	115

GROUP BY

- SELECT **sname**, standing
FROM Students
GROUP BY standing;

sname	standing
???	freshman
???	senior
???	junior
???	sophomore

- Illegal! Can't find one 'sname' for whole table.
- Rule:
 - SELECT fields must either be inside the GROUP BY or aggregates.

GROUP BY

- How about grouping by multiple fields?
- `SELECT standing, gpa, COUNT(*)`

`FROM Students`

`GROUP BY standing, gpa;`

standing	gpa	count
freshman	3	38
freshman	4	21
freshman	2	20
sophomore	4	19
...

HAVING

- HAVING: Modifies which groups are returned.
- SELECT standing, gpa, COUNT(*)

FROM Students

GROUP BY standing, gpa

HAVING COUNT(*) > 30;

standing	gpa	count
freshman	3	38
sophomore	3	32
junior	3	46
junior	4	31

HAVING vs. WHERE

- HAVING: condition on aggregations
 - usually after group by
- WHERE: condition on individual rows
 - usually before group by
- SELECT standing, gpa, COUNT(*)

FROM Students

WHERE sname STARTS_WITH 'A'

GROUP BY standing, gpa

HAVING COUNT(*) > 30;

Logical Order of a Query

- SELECT standing, gpa, COUNT(*)
FROM Students
WHERE sname STARTS_WITH 'A'
GROUP BY standing, gpa
HAVING COUNT(*) > 3;
- Where does aggregation happen?
During GROUP BY!

