

CS186 Discussion Section 7

Aggregation and Query Optimization

10/22/2013

Review: Sorting & Hashing

- a) Consider a simple multi-pass external merge sort algorithm that in the first pass, produces disk-based runs of “B” pages where B is the number of pages of memory available to the algorithm. If B = 50 pages, how many passes (including pass 1) of the algorithm will be required to sort a relation of size 5000 pages and how many runs will be produced by each of the passes?
- b) Suppose we use the in memory sort optimization that generates runs of (on average) length 2B during pass 1. If B = 50, how many passes are required (including pass 1) and how many runs will be produced by each of the passes?
- c) In what scenarios might we prefer to use hashing instead of sorting?

Aggregations/Query Optimization

Consider the following schema:

Car(license, owner_ssn, year, company, model)
Accident(license, accident_date, damage_amount, zipcode)
Owner(ssn, license, name, gender, street, city, zipcode)

NTuples(Car) = 1000 ; NPages(Car) = 100
NTuples(Accident) = 500 ; NPages(Accident) = 20
NTuples(Owner) = 800 ; NPages(Owner) = 50
NDistinct(Car.company) = 50;

For the following questions, assume that we have $B=20$ pages of available in-memory buffer space.

a) For the query: “SELECT gender, COUNT(*) FROM Owner GROUP BY gender;” What is the IO cost for aggregation by sorting?

b) For the query in part a), what is the IO cost for aggregation by hybrid hashing?

c) For the query: “SELECT * FROM Accident A, Car C WHERE A.license = C.license AND A.damage_amount > X;” For what types of values of X would selection push-down significantly improve the cost of the query (Car is the inner table of the join)?”

d) For the query: “SELECT O.name FROM Car C, Owner O WHERE C.license = O.license AND C.company = 'Volvo';” What is the expected cardinality of the Car relation after the initial selections are applied?