

# **CS186 Discussion Section Week 8**

Query Optimization

# Today

1. Review of Optimization Goals
2. Histograms
3. System R (Selinger-style) - optimizer review
4. More time for practice problems

# Review of Optimization Goals

- Why optimize?
  - Better resource utilization.
  - Faster queries. (For some queries, easy to get 99% reduction in query time!)
- What do we optimize?
  - Traditionally - total I/Os +  $f \cdot \text{CPUs}$  to execute query.
  - Could optimize for time to first answer, power consumption, etc.
- What enables optimization?
  - Catalog
  - Relational Algebra
  - Operator model




# Histograms

“.. an estimate of the probability distribution of a continuous variable” - Wikipedia




Often a visual tool - in databases we use them as a concise representation of the distribution of one or more attributes.

Gives us a better estimate of selectivity than “assume things are uniform.”

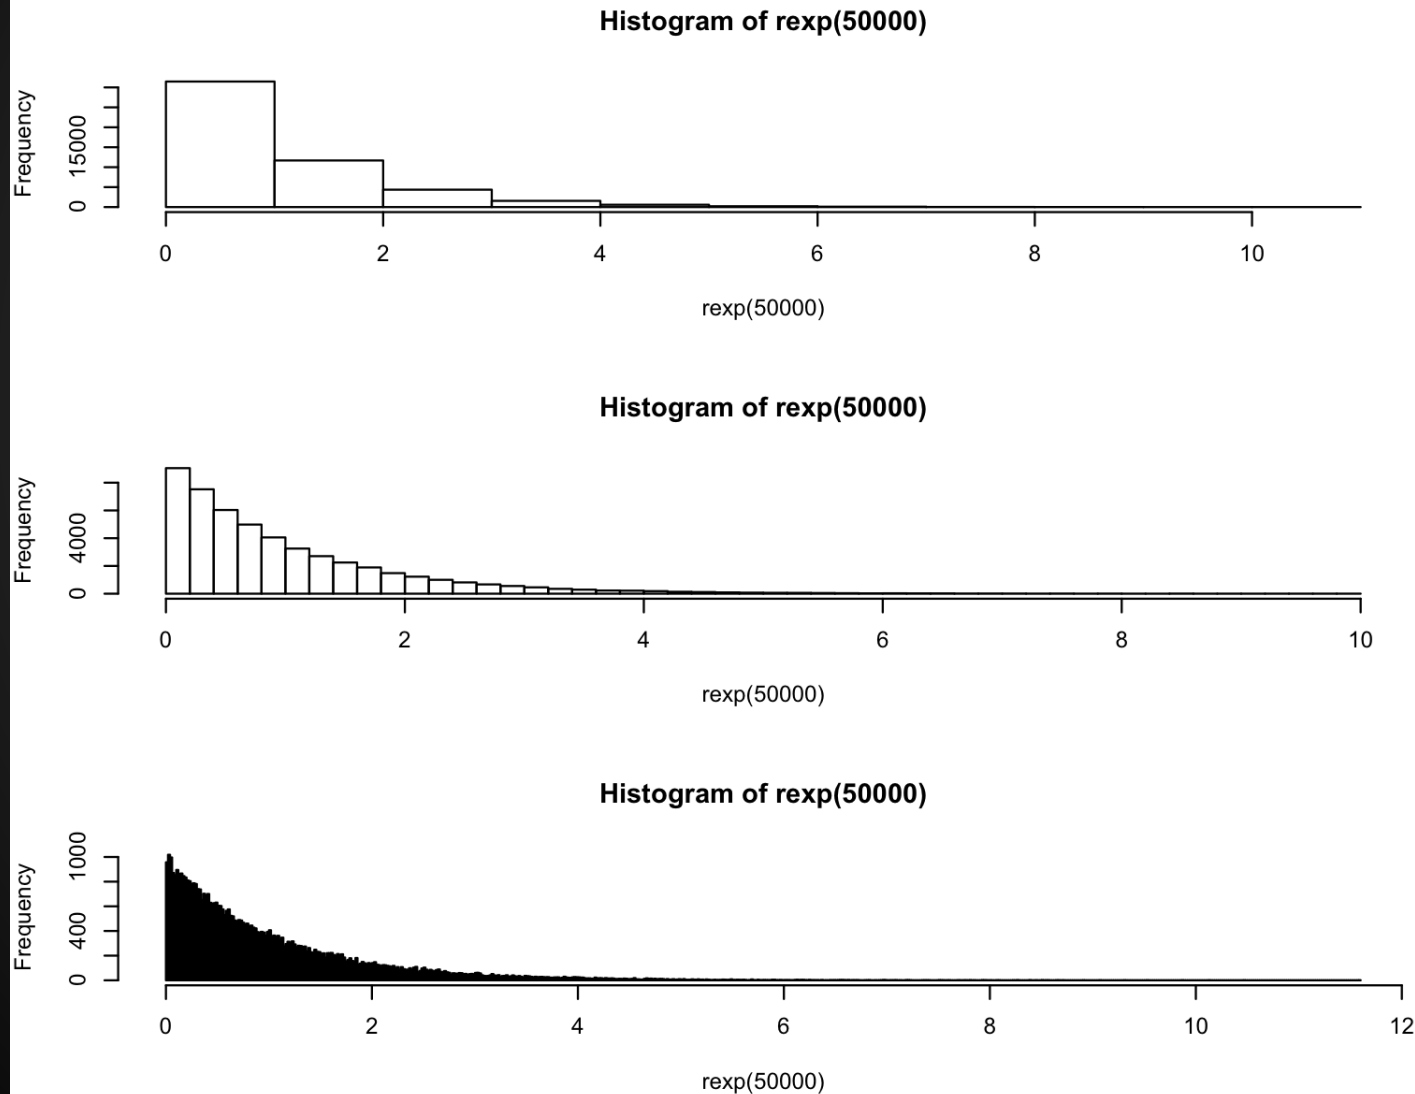
# Histogram Examples - Uniform



# Histogram Examples - Normal



# Histogram Examples - Exponential



# Histogram Wrapup

- Concise representation of data distribution.
- Useful for selectivity estimation.
- Think of it as lossy compression.
- More buckets:
  - better estimates
  - more storage
- You'll implement these!



# Query Optimization - Selinger Style

- System R Optimizer
- Cost model
  - Minimize:  $I/Os + f * CPUs$
- Estimate cost:
  - Estimate cost for each operator - sum them up!
  - Requires selectivity estimates
    - 1/10 if not available!
  - “Interesting Orders” change plan cost
- Prune the search space
  - Left-deep plans only!
- Search the space
  - Dynamic programming!

# Cost Model

A query plan has a single number associated with it - its cost:

$$\text{COST} = \text{IOs} + f * \text{CPUs}$$

One number allows us to say “this plan is better than that plan.”

f - A factor that we can set to determine which is more important - IOs or CPUs.

# Determinants of Plan Cost

- Access method of base tables
  - Scan
  - Index
    - Range vs. lookup
    - Clustered vs. unclustered
- Join ordering
  - Do we want to keep rereading a big table over and over again?
  - What if we have a highly selective scan?
- Join method
  - Sort-merge? Hash? BNL?

# Left Deep Plans

An optimization that accomplishes a few things

1. Prunes search space of possible join orderings from something like  $n!Catalan[n-1] \approx n!4^n$  to something more like  $n!$ 
  - a. [See Catalan numbers.](#)
  - b.  $\#(\text{Left deep plans}) \ll \#(\text{All Plans})$
2. Gives us all “fully pipelined” plans
  - a. Also gives us some plans that are not fully pipelined.

Note: number of possible plans is still  $n!$  !!!

# Dynamic Programming

1. Find the best 1-table access method.
2. Given the best 1-table method as the outer, find the best 2-table.
3. ...
4. Given the best  $(N-1)$ -table method as the outer, find the best  $N$ -table.

\*Wrinkle - instead of “strictly the best” we return the best for each interesting order of the tuples.

\*Wrinkle 2 - do cross products last!

# Interesting Orders

Operator returns an “interesting order” if its result is in order of:

- some **\*ORDER BY\*** attribute
  - means we don’t have to sort later!
- some **\*GROUP BY\*** attribute
  - means we can use the nice scan method for our group-by later!
- some Join attribute of *other* joins
  - Means we can use sort-merge far cheaper!

**Now - practice!**