

CS186 Section

Week 9

Transaction and Concurrency Intro

Today

1. Transactions (ACID)
2. Isolation: Concurrency Control
 - Serializability and conflict serializability
 - 2PL and Strict 2PL
3. Worksheet

Transaction

- Sequence of instructions you want to execute as if one:
 `CookieJar := CookieJar - 1`
 `CookiesAte := CookiesAte + 1`
- Want to ensure:
 - Either both execute, or neither.
 - No other transaction sees only part of this execution:



ACID

- Set of formal properties a transaction has:
 - Atomicity – all or nothing / commit or abort
 - Take cookie and eat it, or don't take cookie
 - Consistency – database remains in consistent state afterwards
 - $\text{CookiesInJar} \geq 0$
 - Isolation – transaction runs as if it's the only one
 - $\text{CookiesInJar} + \text{CookiesAte} == \text{max}$
 - Durability – changes are never lost
 - CookieMonsterApp crashes, wakes back up, CookiesInJar and CookiesAte are whatever they were before.

Isolation - Concurrency Control

- Could just execute 1 transaction at a time...
 - Slow!
- Want to maximize parallelism while maintaining a sense of isolation.
 - “Concurrency Control!”

Serializability

- Serial schedule – “run one at a time”
- Serializable schedule – schedule whose outcomes are equivalent to a serial schedule.

Time	Transaction 1	Transaction 2
1		CookiesBaking := 5
2		Jar := Jar + CookiesBaking
3	Jar := Jar – 1	
4	Ated := Ated + 1	

Serial, serializable, or neither?

Serializability Disclaimer

- **Serializable in databases:**
 - From Wikipedia: “in concurrency control of databases, a transaction schedule is serializable if its outcome is equal to the outcome of its transactions executed serially”
- **Serialization in rest of CS:**
 - In the context of data storage and transmission, serialization is the process of translating data structures or object state into a format that can be stored.

Serializability

- Serial schedule – “run one at a time”
- Serializable schedule – schedule whose outcomes are equivalent to a serial schedule.

Time	Transaction 1	Transaction 2
1		CookiesBaking := 5
2	Jar := Jar – 1	
3	Ated := Ated + 1	
4		Jar := Jar + CookiesBaking

Serial, **serializable**, or neither?


Schedule Abstraction

- Talking about the actual semantics of a program becomes hard
 - instead, we just simplify to Reads and Writes.

Time	Transaction 1	Transaction 2
1		CookiesBaking := 5
2	Jar := Jar - 1	
3	Ated := Ated + 1	
4		Jar := Jar + CookiesBaking

Schedule Abstraction

Time	Transaction 1	Transaction 2
1		CookiesBaking := 5
2	Jar := Jar - 1	
3	Ated := Ated + 1	
4		Jar := Jar + CookiesBaking



Time	Transaction 1	Transaction 2
1		W(CookiesBaking)
2	R(Jar)	
3	W(Jar)	
4	R(Ated)	
5	W(Ated)	
6		R(Jar)
7		W(Jar)

Serializability

- Is this schedule serializable?
 - No!
 - How can we tell in general?

Time	Transaction 1	Transaction 2
1		W(CookiesBaking)
2		R(Jar)
3	R(Jar)	
4	W(Jar)	
5	R(Ated)	
6	W(Ated)	
7		W(Jar)

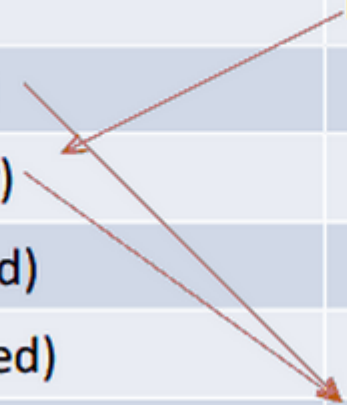
Conflicts

- Conflict:
 - two operations in
 - different transactions on
 - the same object where
 - **at least one is a Write.**


How many conflicts here?

- **Concurrency issues can only arise in the face of conflicts.**

Time	Transaction 1	Transaction 2
1		W(CookiesBaking)
2		R(Jar)
3	R(Jar)	
4	W(Jar)	
5	R(Ated)	
6	W(Ated)	
7		W(Jar)



Time	Transaction 1	Transaction 2
1		W(CookiesBaking)
2		R(Jar)
3	R(Jar)	
4	W(Jar)	
5	R(Ated)	
6	W(Ated)	
7		W(Jar)



Time	Transaction 1	Transaction 2
1		W(CookiesBaking)
2	R(Jar)	
3	W(Jar)	
4	R(Ated)	
5	W(Ated)	
6		R(Jar)
7		W(Jar)

Conflict Serializability

- A schedule is “conflict serializable” if it is equivalent to some **serial schedule with the conflicts in the same order.**
 - A schedule is conflict serializable if and only if its dependency graph is acyclic.
 - Another way to look at it: Sliding operations.

Serializability vs. Conflict Serializability

- Why not just check for serializability?
 - We can verify conflict serializability using dependency graph; much harder to verify serializability.
- Any conflict serializable schedule is serializable.
- Some serializable schedules are not conflict serializable.
- This means we throw out valid schedules!
 - That's ok – it's the price we pay.
- Questions?

Locks

- We use locks to control access to objects.
 - Shared lock: multiple transactions can have a shared lock on the same item (e.g., reading)
 - Exclusive lock: only one (and no other lock) on this item (e.g., writing)

Time	Transaction 1	Transaction 2
1	Lock_X(Jar)	Lock_X(CookiesBaking)
2	R(Jar)	W(CookiesBaking)
3	W(Jar)	Unlock(CookiesBaking)
4	Unlock(Jar)	
5	Lock_X(Ated)	Lock_X(Jar)
6	R(Ated)	R(Jar)
7	W(Ated)	W(Jar)
8	Unlock(Ated)	Unlock(Jar)

Two-Phase Locking (2PL)

- We add a rule for how a transaction may acquire locks:
 - Once you release a lock, you may never acquire a new lock.
- Ensures conflict serializability!
 - Wut, really?
 - In order for a conflict cycle to occur, we need to release a lock so other guy can use our object, then acquire a lock to use the other guy's object

Cascading Aborts and Strict 2PL

- What happens here?

T1: R(A), W(A), **Abort**

T2: R(A), W(A)

- How to fix?
- Strict Two-Phase Locking:
 - All locks held by transaction are only released at the end of the transaction.

Deadlocks

- Dealing with deadlocks:
 - Prevention – stop them from occurring
 - Detection – stop them while occurring
 - In practice: timer

Time	Transaction 1	Transaction 2
1	Lock_X(A) (granted)	
2		Lock_X(B) (granted)
3	Lock_X(B) (waiting)	
4		Lock_X(A) (waiting)
5

Deadlock Prevention

- Disallow deadlocks from ever occurring.
- Two transactions T_{old} and T_{young} .
 - Wait-Die:
 - If T_{old} is waiting for a lock from T_{young} , he just waits.
 - If T_{young} is waiting for a lock from T_{old} , he kills himself.
 - Wound-Wait:
 - If T_{old} is waiting for a lock from T_{young} , he kills T_{young} .
 - If T_{young} is waiting for a lock from T_{old} , he just waits.
- If you die, you restart with ***original*** timestamp.

Deadlock Detection

- Waits-for graph of all transactions.
- If cycle exists, shoot one of the transactions in the cycle.

Summary

- For isolation, we need a serializable schedule.
- Strict 2PL gives us conflict serializability automatically, and avoids cascading aborts.
- Can either detect deadlocks using waits-for graph or prevent it using wait-for or wound-wait

Worksheet!