

# **CS186 Discussion**

## **Section Week 11**

Logging and Recovery

# ACID

- Atomicity
- Consistency
- Isolation
- Durability

# Last week - Isolation

- Isolation: transactions aren't affected by the operations of other transactions.
- We solved this: Strict 2PL gives us serializability

# This week - Atomicity & Durability

- Atomicity: transactions are all or nothing
- Durability: after commit, data never disappears

How can we ensure these if the DB can crash at any time?

# Buffer Policies

- **NO STEAL**: don't let the system “steal” frames with uncommitted updates from the buffer pool and write them to disk
  - Which do we lose if we steal pages with uncommitted data?
    - A. Atomicity
    - B. Durability

# Buffer Policies

- **NO STEAL**: don't let the system “steal” frames with uncommitted updates from the buffer pool and write them to disk
  - Which do we lose if we steal pages with uncommitted data?
    - A. Atomicity
    - B. Durability
  - if there's a crash with STEAL, transaction is incomplete, but some updates are on disk
  - To preserve atomicity, we must UNDO those changes.

# Buffer Policies

- **FORCE**: “force” the buffer manager to write every updated page to disk before committing
  - Which do we lose if we don’t force flushing before committing?
    - A. Atomicity
    - B. Durability

# Buffer Policies

- **FORCE**: “force” the buffer manager to write every updated page to disk before committing
  - Which do we lose if we don’t force flushing before committing?
    - A. Atomicity
    - B. **Durability: commits may not survive a crash**



# Buffer Policies

- **FORCE**: “force” the buffer manager to write every updated page to disk before committing
  - Which do we lose if we don’t force flushing before committing?
    - A. Atomicity: may have written only some of the pages
    - B. Durability: commits may not survive a crash
- if there’s a crash with NO FORCE, we have no guarantee that all committed pages are on disk.
- To preserve atomicity/durability, we must REDO those changes.

# Buffer Policies

- NO STEAL / FORCE (SimpleDB):
  - Don't steal pages with uncommitted data.
  - Always force committed pages to disk.
  - Gives us atomicity and durability! Why not use?
    - Really slow at commit time, especially with both...
- STEAL / NO FORCE (Real World):
  - Can write uncommitted data to disk.
  - Don't have to write data to disk on commit.
  - Essentially, no guarantees about A or D. Why use?
    - Much faster. Can use **write-ahead logging** to get A & D!

# Write-Ahead Logging (WAL)

- Don't guarantee data gets to disk, but guarantee that our **logs** about that data do.
  - Protection from STEAL: Force log record for update out before corresponding data page gets to disk.
  - Protection from NO FORCE: Force all logs for transaction out before commit finishes.
- Log everything!
  - Transaction start
  - Transaction updates
  - Transaction commit
  - Transaction abort

# The Log

- <LSN, pageID, offset, old data, new data, prevLSN>
  - LSN (“Log Sequence Number”): globally increasing ID for log records
  - prevLSN: LSN of the last operation for **this txn**

# More Transaction State

- Transaction Table
  - Answers question: which transactions are currently running?
  - Contains:
    - **XID**: Transaction ID
    - **Status**: Running/Committing/Aborting
    - **lastLSN**: **most recent** LSN created by the txn
- Dirty Page Table
  - Answers question: which buffer pages are dirty?
  - Contains:
    - **pageID**
    - **recLSN**: LSN of **first** update that dirtied this page
- Checkpoints: Occasionally save these tables in the log (helpful if we crash)

# Normal Execution

- Transactions happening, everything bright and cheery.
- Commit occurs: what do we do?
  - Flush the logs to disk!
- Abort occurs: what do we do?
  - We need to undo all the txn's changes.

# Example Log (only 1 txn running)

LSN	Log	prevLSN
10	T1 Start	null

# Example Log (only 1 txn running)

LSN	Log	prevLSN
10	T1 Start	null
20	T1 writes P5	10



# Example Log (only 1 txn running)

LSN	Log	prevLSN
10	T1 Start	null
20	T1 writes P5	10
30	T1 writes P6	20

# Example Log (only 1 txn running)

LSN	Log	prevLSN
10	T1 Start	null
20	T1 writes P5	10
30	T1 writes P6	20
40	T1 Abort	30

# Example Log (only 1 txn running)

LSN	Log	prevLSN
10	T1 Start	null
20	T1 writes P5	10
30	T1 writes P6	20
40	T1 Abort	30
50	CLR: Undo 30, undoNextLSN=20	40

# Example Log (only 1 txn running)

LSN	Log	prevLSN
10	T1 Start	null
20	T1 writes P5	10
30	T1 writes P3	20
40	<span style="color: red;">T1</span> Compensation Log Record	30
50	CLR: Undo 30, undoNextLSN=20	40

# Example Log (only 1 txn running)

LSN	Log	prevLSN
10	T1 Start	null
20	T1 writes P5	10
30	T1 writes P6	20
40	T1 Abort	30
50	CLR: Undo 30, undoNextLSN=20	40
60	CLR: Undo 20, undoNextLSN=null	50

# Example Log (only 1 txn running)

LSN	Log	prevLSN
10	T1 Start	null
20	T1 writes P5	10
30	T1 writes P6	20
40	T1 Abort	30
50	CLR: Undo 30, undoNextLSN=20	40
60	CLR: Undo 20, undoNextLSN=null	50
70	T1 End	60

# Next Week: Crash Recovery

- System crashes. What do we do?
- General plan:
  - Make sure our in-memory txn state is also up to date (analysis)
  - Re-apply changes made by committed txns to make sure they got to disk (needed with NO FORCE)
  - Undo uncommitted changes on disk (needed with STEAL, or for in-flight txns).
- This is called “ARIES”